

Subject: The paper I sent earlier
Date: Wednesday, July 2, 2025 at 12:47:17 PM Central Daylight Time
From: Martin Paulus
To: Dale Peasley, Mikey Ferguson

Hi guys

I have been looking at the paper and had my GPT friend put a more concrete set of steps together how we might use this with our data:

Wednesday, July 2, 2025: Update based on the “*Large Language Models are Powerful Electronic Health Record Encoders*” paper

Below is a scientist-oriented snapshot of the study’s methods and quantitative findings.

- **Purpose** – Test whether general-purpose, instruction-tuned LLM embedding models can encode structured EHRs for clinical prediction, versus an EHR-specific foundation model (CLMBR-T-Base) and a counts + GBM baseline.
- **Datasets**
 - *EHRSHOT* (Stanford): 6,739 patients | 921 k visits | 41.7 M events
 - *UK Biobank* (external): 387,464 patients | 19.5 M visits | 72.3 M events
- **Benchmark tasks** – 15 prediction tasks in four groups (operational outcomes, lab abnormalities, new diagnoses, chest-X-ray findings), evaluated in few-shot regimes (k = 1–128 per class).
- **EHR-to-text serialization** – Structured records converted to Markdown ($\leq 4,096$ tokens): demographics → 24 key LOINC vitals/labs → visit summary → reverse-chronological event lists with ontology-based code descriptions.
- **Embedding models & training**
 - LLM encoders: **GTE-Qwen2-7B-Instruct** & **LLM2Vec-Llama-3.1-8B-Instruct** generate patient embeddings; logistic-regression heads used for all tasks.
 - Baselines: CLMBR-T-Base embeddings and counts + GBM evaluated in the same pipeline.
- **EHRSHOT performance (macro-AUROC, 95 % CI)**

Model	Macro-AUROC	With CLMBR concat
GTE-Qwen2-7B	0.755 (0.730–0.780)	0.786 (0.762–0.811)
LLM2Vec-Llama-3.1-8B	0.734 (0.707–0.762)	0.773 (0.747–0.799)
CLMBR-T-Base	0.769 (0.746–0.792)	—
Counts + GBM	0.719 (0.691–0.748)	—
- **UK Biobank external validation (macro-AUROC)**
 - | Task group | GTE-Qwen2-7B | CLMBR-T-Base | Llama-8B |
 - | Mortality | **0.826** | 0.782 | 0.781
 - | Hospitalization | **0.655** | 0.639 | 0.638
 - | 23-disease onset | **0.727** | 0.707 | 0.715
 - | Macro avg. | **0.736** | 0.709 | 0.711
- **Few-shot behavior** – LLM embeddings yield the largest gains over CLMBR and GBM when 4–32 labeled examples are available; benefits taper as data scale increases.
- **Context & recency effects**

- Best macro-AUROC for GTE at 4,096 tokens; Llama peaks at 2,048 tokens, with no gain from 8,192 tokens.
- Limiting input to the most recent **1 month** of history performs on par with full histories, hinting at stronger predictive signal in recent events.
- **Limitations flagged by authors** – Hand-crafted serialization may bias comparisons; 4 k-token cap can omit long-term context; GPU cost high (≈ 20 h on $8 \times A100$ for GTE; ≈ 30 h for Llama); results drawn from only two health-system datasets.

Step-by-step blueprint for turning structured EHRs into GPT-ready inputs and predictive features

- 1. Anchor the clinical question up front**
Pick a reference (“prediction-time”) date and define the outcome window—e.g., one-year mortality or 30-day readmission. Normalizing every event timestamp to that single date simplifies temporality downstream.
- 2. Extract & harmonize raw codes**
Pull all structured tables (diagnoses, meds, labs, procedures, vitals) and map every code to a common “ontology/code” string. Resolve code \rightarrow description via prepared ontologies (SNOMED, RxNorm, LOINC, CPT-4, ICD-10-PCS, etc.) and add custom maps for gaps.
- 3. Serialize the record into Markdown**
Use a lightweight, strictly ordered text template because Markdown “adds minimal overhead” yet preserves structure for the model.
 1. # Electronic Healthcare Record
 2. Current time: [YYYY-MM-DD]
 3. ## Patient Demographics
 4. ...
- 4. Populate the template (information-rich \rightarrow general)**
 - 4.1 **Demographics first** (age, sex, ethnicity). Convert DOB to age.
 - 4.2 **Recent high-value time series**: choose ~ 24 canonical vitals/labs, keep the last three values, add units plus low/normal/high flags.
 - 4.3 **Visit summaries**: one-line bullets of each visit to guard against later truncation.
 - 4.4 **Detailed reverse-chronological events** (conditions \triangleright meds \triangleright procedures).
- 5. Cap length & control recency**
Limit the serialized text to the context window you can afford (e.g., 4 096 tokens). If GPU memory allows, test longer windows (8 192) or chunk averaging; in practice, a **one-month** window often equals full-history performance.
- 6. (Optional) Prepend task-specific instructions**
Short natural-language prompts (“Predict whether the patient will be hospitalized in the next year”) can boost embeddings—verify via ablation by removing or generalizing instructions.
- 7. Choose an LLM embedding backbone**
Select a decoder-only model already converted to an embedding variant (e.g., **GTE-Qwen2-7B-Instruct** or **LLM2Vec-Llama-3.1-8B**) that can ingest ≥ 4 k tokens.
Feed the serialized text; mean-pool the last hidden layer to obtain a fixed-length vector.
- 8. Attach a lightweight predictor**
Logistic regression or gradient-boosted trees on top of the frozen embeddings give calibrated probabilities while keeping training cheap.
- 9. Evaluate systematically**
Follow a few-shot protocol ($k = 1 \rightarrow 128$ positives + negatives) and report AUROC/AUPRC with

bootstrapped 95 % CIs.

10. **Run ablations & sensitivity analyses**

Remove demographics, aggregated labs, visit summaries, or instructions one at a time to see which components drive performance.

11. **Plan for production constraints**

Expect ~20 h on 8×A100 GPUs for 6 k patients at 4 k tokens with a 7-B model; budget >30 h for an 8-B model or a 400 k-patient external cohort.

12. **Document limitations & future refinements**

Manual serialization can bias results; token limits may drop older context; large models add latency and cost. Consider serialization-free approaches or model distillation for deployment.

In practice: start with Steps 1–5 to create a reproducible, compact Markdown view of each patient; plug that into Step 7’s embedding model; and iterate through Steps 8–12 until the predictive metrics and runtime satisfy your clinical use-case.

Critical implementation issues highlighted by the manuscript

Domain	What to watch for	Evidence from the paper
1. Data preprocessing & serialization	<ul style="list-style-type: none">• <i>Subjective template design:</i> manually deciding which tables, time windows, and concepts go first can bias the model toward the target tasks.• <i>Ontology mapping headaches:</i> every OMOP/SNOMED/ICD/RxNorm/LOINC code must be mapped to a human-readable string—rare or site-specific codes easily slip through the cracks.• <i>Signal-to-noise filtering:</i> the authors had to merge synonyms, pick 24 “canonical” vitals/labs, and clip implausible lab values. Skipping a similar quality-control pass will flood the tokenizer with junk tokens.	
2. Context length & temporal coverage	<ul style="list-style-type: none">• A hard 4 096-token cap means older visits may be truncated.• Ablations show that one-month of history often matches or beats the full record—so blindly serializing the entire chart wastes tokens and GPU.	
3. Instruction engineering	<ul style="list-style-type: none">• Task-specific prompts materially change embedding quality; they also inflate the workload because each (patient × task) pair must be re-encoded.• Performance is sensitive to the <i>exact wording</i> of those prompts, which hurts reproducibility across sites or versions.	
4. Model architecture & downstream heads	<ul style="list-style-type: none">• LLM embeddings require a separate classifier (logistic regression in the study); you lose the native zero-shot capacity that makes GPT attractive.• Embeddings and CLMBR vectors are complementary—concatenating them gave the single best AUROC. Plan for model ensembling if you want peak accuracy.	
5. Compute &	<ul style="list-style-type: none">• Encoding 6 739 patients (EHRSHOT) at 4 096 tokens	

latency costs

needed ≈ 20 h on $8 \times \text{A100 GPUs}$ for a 7-B model; scaling to 387 k UK Biobank patients took ≈ 35 h per task.

- Larger models (8 B) hit memory errors unless you batch-split, further slowing the pipeline.

6. Generalization limits

- The study tested only two datasets; differences in coding practices or data density at other hospitals may erode gains.
- CLMBR's drop on UK Biobank underscores how *unseen codes* hurt specialized models—LLM approaches mitigate but don't eliminate this risk.

7. Bias, fairness & calibration

- Manual choices about what to include first (e.g., labs vs. medications) can privilege some conditions over others.
- LLMs are agnostic to coding systems, but their embeddings still need calibration layers; otherwise probability outputs are hard to trust clinically.

Practical take-aways

1. **Design serialization before training:** lock the template, code maps, and time window, then version-control it to curb “template drift.”
2. **Pilot recency windows:** start with 30 days and expand only if specific tasks demand long-term context.
3. **Budget GPUs realistically:** a single full-scale encode pass can be a multi-day cluster job—plan queuing and fault-tolerance.
4. **Keep a calibration layer:** temperature scale or isotonic-regress the logistic head so bedside probabilities remain reliable.
5. **Document every prompt:** treat instruction texts as hyper-parameters; store them alongside model weights for auditability.

Addressing these points early will spare costly re-runs and smooth the path from proof-of-concept to a deployable clinical prediction service.

Below is a concrete, end-to-end blueprint for turning structured EHR data into a single GPT-friendly text string that can be embedded and used to predict **12-month depression deterioration** (e.g., PHQ-9 ≥ 15 or new MDD hospitalization). It is written as if you were about to implement it tomorrow on a real data warehouse; feel free to copy-paste the code blocks as starting templates.

1 Fix the prediction target and index time

index_date = visit_date # index = most recent outpatient encounter

outcome_horizon = 365 # days to look forward

label = 1 if phq9 \geq 15 or mdd_hosp within 365d else 0

Rationale – anchoring every record to a single date lets you standardize “past-to-future” directionality and build rolling cohorts easily.

2 Pull & harmonize raw tables (OMOP style)

Domain	Key fields kept	Extras
Condition	condition_concept_id, start_date	map ICD-10 F32., F33. \rightarrow “MDD”
Drug exposure	drug_concept_id, start_date, days_supply	SSRIs, SNRIs, antipsychotics
Measurement	loinc_concept_id, value, unit, date	TSH, HbA1c, CRP, vit D, LDL
Observation	observation_concept_id, value	PHQ-9, GAD-7, BMI, smoking status

Visit

visit_concept_id, date, type

“Psychiatry-OP”, “ED”, “Inpatient”

Convert concept_ids to **short human-readable tokens** (one time only).

Example:

```
SELECT concept_id, LOWER(REPLACE(concept_name, ' ', '_')) AS token
INTO concept_tokens
FROM vocabulary.concept
```

3 Shallow feature engineering

1. **Temporal bucketing** – tag each event with $\Delta t = \text{index_date} - \text{event_date}$ in days.
2. **Value flagging** – for every continuous measure create discrete flags: low/normal/high.
3. **Medication adherence surrogate** – rolling 90-day medication possession ratio (MPR) for antidepressants.

4 Serialize into deterministic Markdown ($\leq 4\,096$ tokens)

```
# ELECTRONIC_HEALTH_RECORD
current_time: 2025-03-14
prediction_window: next_365_days
## PATIENT
age: 46
sex: female
race: white
## RECENT_METRICS
bmi: 31 (high)
phq9: 12 (moderate)
gad7: 7 (mild)
tsh: 3.8 µIU/mL (normal)
## MED_HISTORY_LAST_365_DAYS
[-10d] condition: major_depressive_disorder (F33.1)
[-10d] drug_start: sertraline 50 mg qd (days_supply=30)
[-45d] measurement: crp 6 mg/L (high)
[-60d] visit: psychiatry_outpatient
[-120d] measurement: hba1c 7.2 % (high)
[-200d] drug_start: bupropion_sr 150 mg bid
## MED_HISTORY_-366_to_-1095_DAYS
... (truncate middle events if token budget exceeded)
## SOCIAL
smoking: former
alcohol: occasional
exercise_level: low
```

Determinism is crucial – the same patient on the same date always yields identical text.

5 Embed once per patient

```
from transformers import AutoTokenizer, AutoModel
tok = AutoTokenizer.from_pretrained("thenlper/gte-small") # 8 k ctx
model = AutoModel.from_pretrained("thenlper/gte-small")
tokens = tok(ehr_markdown, truncation=True, padding=False, return_tensors='pt')
emb = model(**tokens).last_hidden_state.mean(dim=1) # shape (1, 768)
```

Tip – cache the embedding table; downstream experiments become a logistic-regression fit, not a

GPU re-encode.

6 Lightweight predictor

```
from sklearn.linear_model import LogisticRegressionCV
clf = LogisticRegressionCV(Cs=10, cv=5, max_iter=5000, class_weight='balanced')
clf.fit(embeddings_train, y_train)
y_hat = clf.predict_proba(embeddings_test)[: ,1]
```

Evaluate AUROC/AUPRC and calibration (e.g., Hosmer–Lemeshow, ECE).

Keep a temperature-scaling layer if probabilities will drive clinical alerts.

7 Ablation tests (one line each)

```
python run.py --drop-section RECENT_METRICS
python run.py --window 90 # only 3 mo history
python run.py --remove-labs
python run.py --strip-meds
```

Track which chunks move performance the most; typical finding: **recent PHQ-9 trajectory + antidepressant MPR** dominates.

8 Deployment checklist

Item	Why it matters	Quick mitigation
Token cap	4096-token GPT-4o window \approx 25 k chars	enforce size guard, drop oldest events
PHI risk	Markdown is still PHI	encrypt at rest, purge names/addresses
Latency	150 ms per patient (7-B on A100, fp16)	nightly batch encode; serve cached vector
Drift	Coding practice changes (ICD-11, new meds)	weekly concept-map diff + re-train trigger

TL;DR

1. **Anchor** each chart at the prediction visit.
2. **Map** all codes to short tokens.
3. **Serialize** in a fixed Markdown scaffold (NEWEST \rightarrow OLDEST).
4. **Embed** with a 7-B instruction-tuned model once.
5. **Fit** a simple classifier; iterate with ablations.

With one evening of SQL + Python you’ll have a solid, auditable depression-risk encoder that can be swapped into any GPT-friendly pipeline.

Turning the GPT-based EHR encoder into a digital-twin factory

Below is a concrete technical pathway that starts with the Markdown-serialization/LLM-embedding pipeline we just designed for 12-month depression risk and extends it into **individualized, counterfactual simulators**—i.e., digital twins that let you ask “*what if I adjust medication X, start CBT today, or change BMI by 3 kg/m²—how does that alter this patient’s future PHQ-9 or hospitalization risk?*”

Stage	Objective	Design choices & tips
1. Canonical state vector	Freeze the baseline encoder you already validated (Steps 1–5 earlier). Each patient-timepoint is a 768-D vector st capturing <i>current</i>	• Embed <i>every</i> monthly snapshot for the past N months (or per visit).

	demographics, labs, meds, diagnoses, recent PHQ-9 trajectory, etc.	<ul style="list-style-type: none"> • Store (patient_id, t, st) in a feature store so you can query history fast.
2. Action vocabulary	Define a compact, machine-readable set of <i>interventions</i> at : SSRIs \uparrow/\downarrow , start CBT, add bupropion, \uparrow exercise flag, BMI-3 %, etc.	<ul style="list-style-type: none"> • Keep ≤ 50 discrete actions to avoid data sparsity. • Encode “no change” as an explicit action.
3. Dynamics model (st, at) \rightarrow st+Δ	Learn how state evolves under actions. Two viable architectures: (a) GPT-style autoregressive transformer that inputs <i>serialized Markdown</i> + tokens and predicts next-month Markdown, then re-encode. (b) Lightweight MLP/RNN that takes $(st \oplus at)$ and outputs Δ s. Pick (b) if speed matters; (a) if richness of generated notes is critical.	<ul style="list-style-type: none"> • Train with teacher forcing on real sequences: minimize MSE for continuous labs, CE for discrete codes, KL for embedding drift. • Include a <i>calendar offset</i> embedding so the model knows Δ days.
4. Outcome heads	Attach differentiable heads to s that predict clinical endpoints (PHQ-9, inpatient MDD stay, suicidality flag). You already have a logistic-regression head; keep it frozen or co-train.	<ul style="list-style-type: none"> • Multi-task learning stabilizes the dynamics model—every future state is supervised by all available outcomes.
5. Twin generator	For a real patient at $t = 0$: 1. Encode to s0 . 2. Choose counter-factual action sequence {$\hat{a}0:T-1$} (e.g., “increase SSRI dose at $t = 0$, add CBT at $t = 1$ ”). 3. Roll the dynamics model forward T steps \rightarrow $\hat{s}1:T$. 4. Pipe each $\hat{s}t$ through the outcome heads to yield risk trajectories.	<ul style="list-style-type: none"> • Sampling: add Gaussian noise to latent Δs or use nucleus sampling on the autoregressive model to capture uncertainty. • Generate $N = 100$ twins per strategy and average risks to estimate credible intervals.
6. Calibration & realism checks	Internal validity : compare simulated trajectories with held-out real patients who actually received those actions. Face validity : have clinicians judge generated Markdown snippets (if using the autoregressive path).	<ul style="list-style-type: none"> • Use <i>energy distance</i> or <i>K-nearest-neighbor MMD</i> between real and synthetic sequences. • Track coverage of lab trajectories (e.g., HbA1c cannot jump 5 % in a month).
7. Policy evaluation loop	Wrap the twin generator in an API: <code>get_twin(patient_id, action_plan_json, horizon=365)</code> \rightarrow JSON of monthly outcome probabilities & synthetic notes.	<ul style="list-style-type: none"> • Cache embeddings; only the dynamics roll-out needs GPU (often feasible on a single A10). • Add a Shapley or Integrated-Gradients explainer on $\hat{s}t$ to surface <i>why</i> risk changed.
8.	Digital twins drive decision support, so document assumptions & biases: coding drift,	<ul style="list-style-type: none"> • Maintain model cards with: data window, action

Minimal working code sketch (MLP dynamics)

```
# s_dim = 768, a_dim = 50 one-hot, hidden = 1024
dyn = torch.nn.Sequential(
    torch.nn.Linear(768 + 50, 1024), torch.nn.ReLU(),
    torch.nn.Linear(1024, 768)
)

for (s_t, a_t, s_next) in dataloader:
    pred = s_t + dyn(torch.cat([s_t, a_t], dim=-1)) # residual update
    loss = torch.nn.MSELoss()(pred, s_next)
    loss.backward(); opt.step()

Generate a twin:
s = s_0.clone()
traj, risks = [], []
for a in planned_actions:      # list of one-hot vectors
    s = s + dyn(torch.cat([s, a], -1))
    traj.append(s)
    risks.append(outcome_head(s).sigmoid())
```

Why this qualifies as a digital twin

1. **Individualized** – seeded with the patient’s own encoded EHR state.
2. **Counter-factual** – you can swap in any plausible action sequence.
3. **Dynamic** – generates month-by-month physiological and diagnostic evolution, not just a static risk score.
4. **Interpretable** – outcome heads plus attribution on latent dimensions show which features drive divergence.

With this pipeline, a clinician can interactively explore “*What if we augment sertraline with bupropion?*” and immediately see projected PHQ-9 trajectories, hospitalization risk, and even synthetic progress-note snippets—all while the underlying state vectors remain de-identified and compact enough for real-time use.

The Art of War: In the midst of chaos, there is also opportunity.
Martin

Scientific Director and President, Laureate Institute for Brain Research
6655 S Yale Ave Tulsa, OK 74136-3326 P 918 502 5120 F 918 502 5135
email: mpaulus@laureateinstitute.org X: @mpwmpaulus
web: <http://www.laureateinstitute.org>

Professor of Neuroscience, Oxley College of Health and Natural Sciences
University of Tulsa, email: mpp4692@utulsa.edu

Adjunct Professor, Department of Psychiatry, University of California, San Diego
email: mpaulus@health.ucsd.edu